

SproutJourney: Reinventing Early Education in India with AI

Problem Statement

In India, children between ages 6–12 often fall through the cracks of a rigid, outdated education system that emphasizes rote memorization over cognitive development. Many, like Aarav—a real child with learning disabilities—find classroom learning disheartening due to a lack of personalization. The system fails to adapt to how children actually learn. Without intervention, early disengagement cascades into long-term academic and social challenges. There is a critical need for a learning platform that evolves with each child—understanding not only what they know but *how* they understand it—especially in foundational years where curiosity, ethics, and critical thinking are most malleable.

Target Audience & Context

SproutJourney is built for:

- **Primary school children (ages 6–12)**, especially those who are neurodiverse, in underserved regions, or disengaged from traditional systems
- **Parents and educators** seeking inclusive, scalable, personalized learning solutions

India's education system, serving 260M+ students, is predominantly rote-based. Regional disparities and teacher shortages exacerbate this. Children learn through different cognitive routes—visual, narrative, logical—but the system offers only one-size-fits-all teaching. SproutJourney uses AI to bring inclusivity, adaptability, and emotional intelligence to the classroom.

Use of Gen-AI

SproutJourney is powered by a **custom fine-tuned LLM based on Gemma**, layered with a **multi-character multi-agent system**, and integrated with a **Retrieval-Augmented Generation (RAG) pipeline** grounded in a **custom language graph**.

Each subject-specific AI character is trained on:

- NCERT-aligned curriculum
- Culturally rooted content (e.g., Bhagavad Gita, Indian epics)
- Application-based and rote-style educational prompts curated with a child-safe filter

The language graph serves as a concept map that connects grade-wise knowledge clusters to storytelling, ethical reasoning, and pedagogical prompts. Each agent accesses domain-specific RAG nodes with **guardrails** to ensure content safety—proactively filtering out inappropriate responses and hallucinations.

The AI learns in real time, updating a child's **cognitive fingerprint** based on interaction style, retention rates, and emotional engagement. This enables:

- Precision adaptation to learner type (visual, conversational, gamified)
- Safe and relevant content flow, even in edge-case queries
- Meaningful inclusion for children with learning disabilities, language barriers, or early-stage reading abilities

Solution Framework

SproutJourney's technical core centers around a **child-safe, emotionally aware AI learning engine** with structured personalization.

System Architecture:

- **Gemma-based LLM** fine-tuned using LoRA and child-safe corpora
- **Agent mesh network**: Each character (e.g., Eli, Nyra, Viv) has its own instruction set, vocabulary range, and tone model for its domain
- **Domain-specific RAG modules**: Each subject pulls content from a vector store linked to a grade-wise **language graph**
- **Language graph**: A dynamic knowledge topology mapping concepts, Bloom's taxonomy levels, and their moral or real-world connections. This powers the storytelling logic and follow-up questioning
- **Safety-first guardrails**: Filters are applied at input pre-processing and output stages to prevent unsafe content, guided by pre-labeled examples and a custom moderation layer

Tracking & Personalization:

- **Cognitive fingerprint engine**: Continuously updates student profile using NLP and interaction analytics
- Flags patterns like attention lapses, misunderstanding zones, and emotional tone via response vectors

Learning Modes:

1. **StoryTime** – agent-led cultural narratives grounded in cognitive themes
2. **LearningTime** – visual or audio explanation, Socratic method-based
3. **ThinkingTime** – applied reasoning via moral dilemmas or real-life math/finance problems
4. **FunTime** – gamified reinforcement with narrative callbacks

We are currently developing dedicated tracks for **financial literacy**, **civic responsibility**, and **practical ethics**, designed to prepare learners for real-world decision-making. We are also experimenting with **video generation AI**, **image generation AI** (e.g., Stable Diffusion models), and **TTS (text-to-speech)** using **Eleven Labs** and **Azure AI Voice**, enhancing engagement especially for low-literacy or visually-oriented learners. We have implemented

LangChain and **LangGraph** for orchestration, and are currently testing with **ChromaDB**, with plans to migrate to a **scalable vector database** for production.

Feasibility & Execution

The system is built using Unity (front-end), Gemma with LoRA fine-tuning (LLM), and a modular Python-based orchestration layer utilizing LangChain and LangGraph for flow control. Domain-specific RAGs are currently served via **ChromaDB**, with plans to shift to a more scalable vector database. Language graph generation was done with spaCy + NetworkX, and child-safe datasets were curated from NCERT, filtered question banks, and application-based content. Safety and hallucination control are enforced through input preprocessing filters and output moderation. AI-generated visuals and audio are being prototyped using Stable Diffusion, Eleven Labs, and Azure AI voices.

Scalability & Impact

SproutJourney's architecture supports horizontal scaling by agent (subject/domain) and vertical scaling by age/grade. The modular vector stores and RAG graphs allow region-specific language and topic packs. With real-time adaptive profiling and low-compute edge deployment support, SproutJourney can work even in low-connectivity classrooms. Inclusivity extends from early learners to children with dyslexia or ADHD. In the long term, the platform enables longitudinal mapping of thought patterns, not just learning scores—turning education into a personalized growth journey. Its safety-first design positions it as a reliable platform for national-scale foundational education transformation.

48-Hour Hackathon Plan

- Hours 0–12: Set up LangGraph orchestration and host fine-tuned Gemma model on the cloud; implement basic agent interactions and RAG pipeline.
- Hours 12–36: Build Unity-based game interface with key learning modes. Integrate TTS and image generation. Prepare subject-wise cognitive agent flows with character voice assignment.
- Hours 36–48: Finalize cognitive fingerprinting engine, add real-time personalization logic. Prepare pitch materials, demo walkthroughs, and documentation for submission.

Summary & Minimum Lovable Product

SproutJourney is not just another AI tutor—it is an emotionally intelligent, culturally grounded, and technically robust education engine. Our MLP includes a working system of AI agents with subject-specific personality, cognitive tracking, and safety guardrails. Designed for real-world inclusivity and pedagogical integrity, it evolves as a companion for the child, not a content pusher.